

# How to Properly Blame Things for Causing Latency

An introduction to Distributed Tracing and Zipkin

**Adrian Cole, Pivotal**  
@adrianfcole

# Introduction

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up

@adriancole

spring cloud at pivotal  
focused on distributed tracing  
helped open zipkin

# Understanding Latency

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up

# Understanding our architecture

Microservice and data pipeline **architectures are often a graph of components**, distributed across a network.

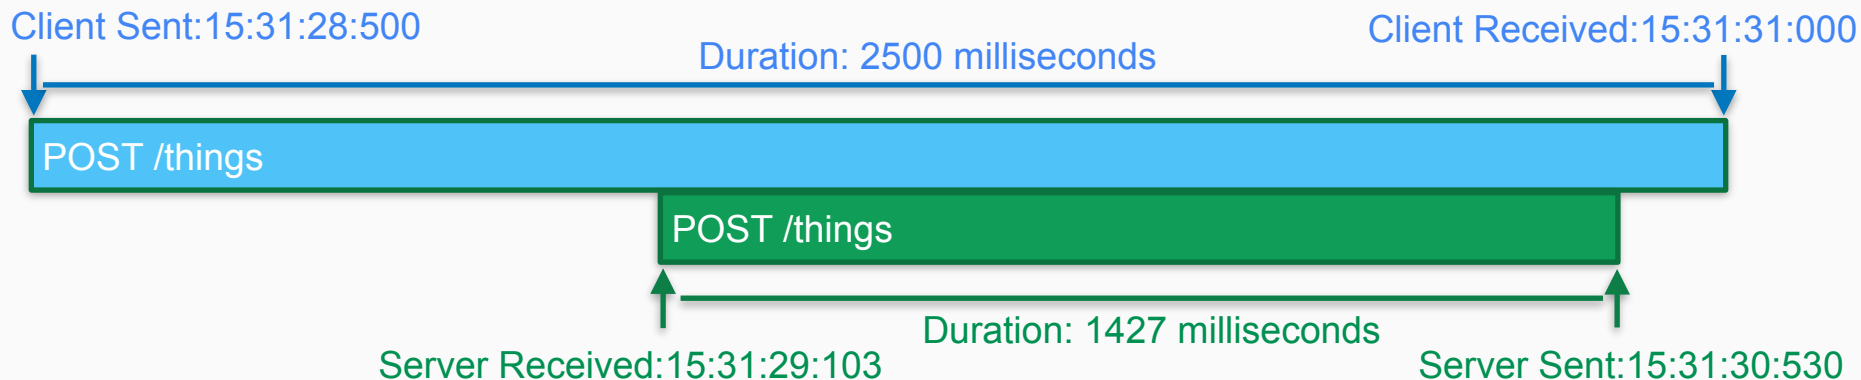
A **call graph or data flow can become delayed** or fail due to the nature of the operation, components, or edges between them.

We want to **understand our current architecture** and troubleshoot latency problems, **in production**.

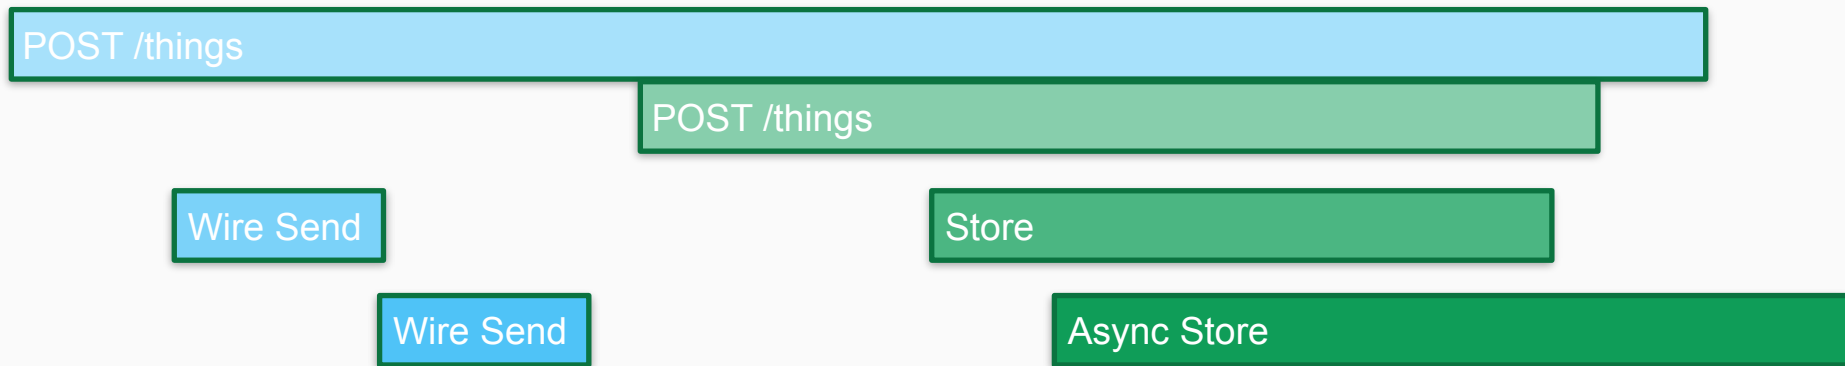
# Why is POST /things slow?

POST /things

# There's often two sides to the story

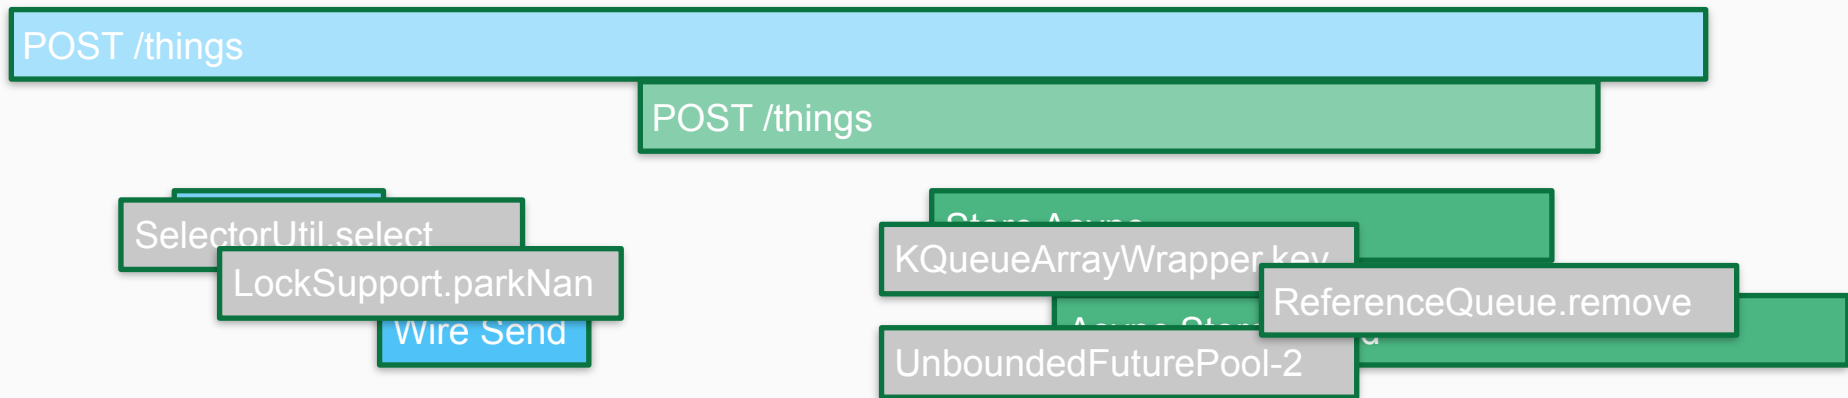


and not all operations are on the critical path






# and not all operations are relevant



# Service architecture isn't this simple anymore


Single-server scenarios aren't realistic or don't fully explain latency.



# Can we make troubleshooting wizard-free?

We no longer need wizards to deploy complex architectures.

We shouldn't need wizards to troubleshoot them, either!



# Distributed Tracing

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up

# Distributed Tracing commoditizes knowledge

Distributed tracing systems collect end-to-end latency graphs (traces) in near real-time.

You can compare traces to understand why certain requests take longer than others.

# Distributed Tracing Vocabulary

A **Span** is an individual operation that took place. A span contains **timestamped events** and **tags**.

A **Trace** is an end-to-end latency graph, composed of spans.

# A Span is an individual operation

Operation

POST /things

wombats:10.2.3.47:8080

Events

Server Received

Server Sent

Tags

peer.ipv4	1.2.3.4
http.request-id	abcd-ffe
http.request.size	15 MiB
http.url	...&features=HD-uploads

# Tracing Systems are Observability Tools

Tracing systems collect, process and present data reported by tracers.

- aggregate spans into trace trees
- provide query and visualization focused on latency
- have retention policy (usually days)



# ProTip: Tracing is not just for latency

Some wins unrelated to latency

- Understand your architecture
- Find services that aren't used
- Reduce time spent on triage

# Zipkin

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up

# Zipkin is a distributed tracing system

Duration: 209.323ms

Services: 5

Depth: 7

Total Spans: 24

JSON

Expand All

Collapse All

Filter Service Se... ▼

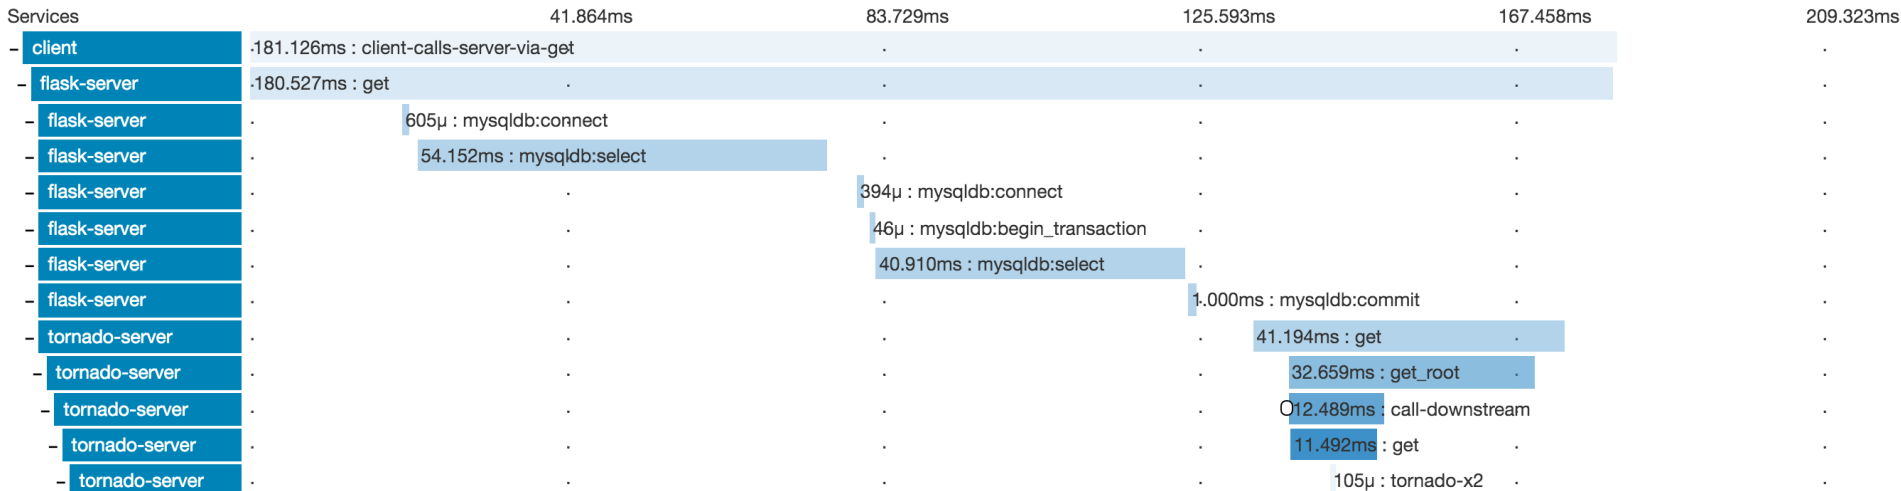
client x4

flask-server x10

missing-service-name x2

tchannel-server x2

tornado-server x11



# Zipkin lives in GitHub

Zipkin was created by Twitter in 2012. In 2015, OpenZipkin became the primary fork.

OpenZipkin is an org on GitHub. It contains tracers, OpenApi spec, service components and docker images.

<https://github.com/openzipkin>

# Zipkin Architecture

Tracers **report** spans HTTP or Kafka.

Servers **collect** spans, storing them in MySQL, Cassandra, or Elasticsearch.

Users **query** for traces via Zipkin's Web UI or Api.

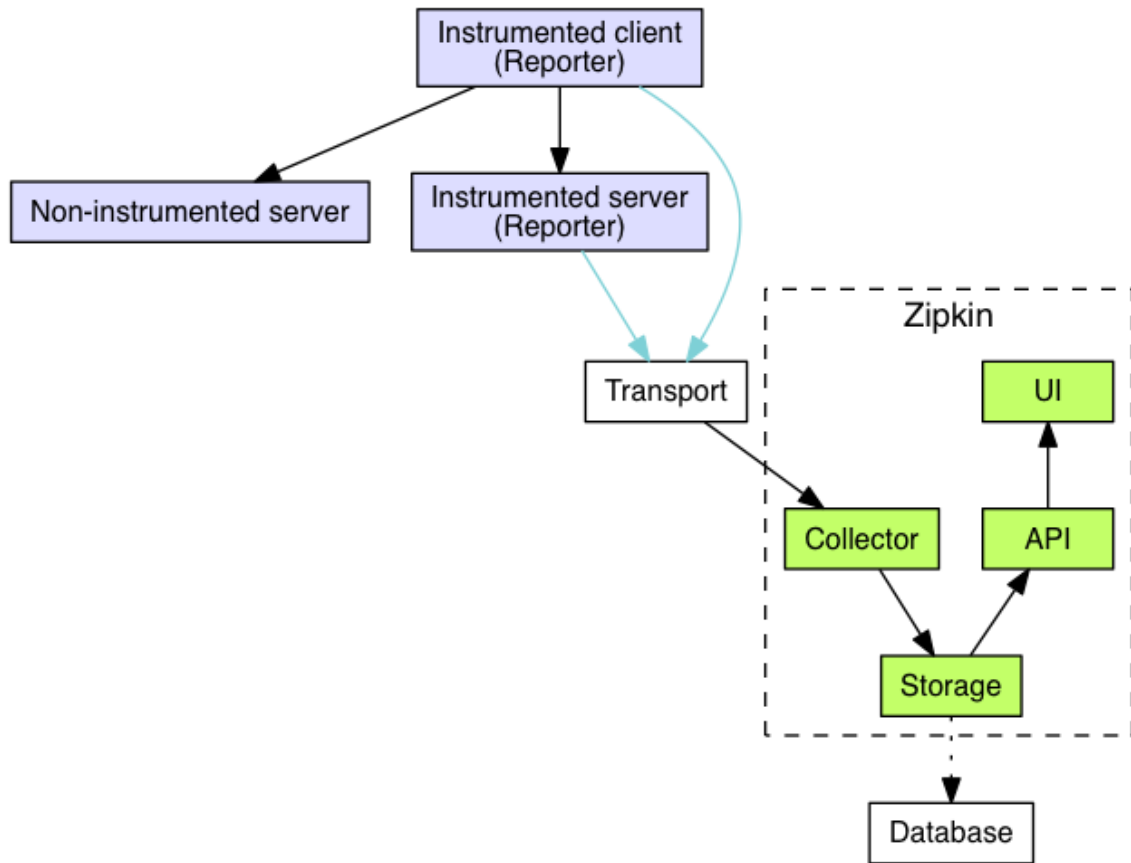
Platform frameworks for Zipkin:

[Bosh \(Cloud Foundry\)](#)

[Docker](#) (in Zipkin's org)

[Kubernetes](#)

[Mesos](#)



# Zipkin has starter architecture

Tracing is new for a lot of folks.

For many, the MySQL option is a good start, as it is familiar.

```
services:
  storage:
    image: openzipkin/zipkin-mysql
    container_name: mysql
    ports:
      - 3306:3306
  server:
    image: openzipkin/zipkin
    environment:
      - STORAGE_TYPE=mysql
      - MYSQL_HOST=mysql
    ports:
      - 9411:9411
    depends_on:
      - storage
```

# Zipkin can be as simple as a single file

```
$ curl -SL 'https://search.maven.org/remote_content?g=io.zipkin.java&a=zipkin-server&v=LATEST&c=exec' > zipkin.jar
$ SELF_TRACING_ENABLED=true java -jar zipkin.jar
```

[illegible]

```
2016-08-01 18:50:07.098 INFO 8526 --- [main] zipkin.server.ZipkinServer : Starting ZipkinServer on acole
with PID 8526 (/Users/acole/oss/sleuth-webmvc-example/zipkin.jar started by acole in /Users/acole/oss/sleuth-webmvc-example)
-snip-
```

```
$ curl -s localhost:9411/api/v1/services|jq .
[
  "zipkin-server"
]
```

# Demo

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up



# Distributed Tracing across Spring Boot apps

Two Spring Boot (Java) services collaborate over http.

Zipkin will show how long the whole operation took, as well how much time was spent in each service.

<https://github.com/openzipkin/sleuth-webmvc-example>

# Spring Cloud Sleuth

Java

Web requests in the demo are served by Spring MVC controllers. Tracing of these are automatically performed by Spring Cloud Sleuth.

Spring Cloud Sleuth reports to Zipkin via HTTP by depending on spring-cloud-sleuth-zipkin.

<https://cloud.spring.io/spring-cloud-sleuth/>

# Wrapping Up

introduction
understanding latency
distributed tracing
zipkin
demo
wrapping up

# Wrapping up

Start by sending traces directly to a zipkin server.

Grow into fanciness as you need it: sampling, streaming, etc

Remember you are not alone!

[gitter.im/openzipkin/zipkin](https://gitter.im/openzipkin/zipkin)

[gitter.im/spring-cloud/spring-cloud-sleuth](https://gitter.im/spring-cloud/spring-cloud-sleuth)